

DocMind AI -- Research Paper Benchmark

Document Ingestion & Multi-Head Attention Verification

Page 1 Overview

This sample document demonstrates grounded vector retrieval and citation jumping.

DocMind AI processes multi-page PDFs with real layout rendering and page-level evidence highlighting.

Page 2 -- Empirical Metrics

ResNet-101 model achieved 96.1% accuracy on ImageNet dataset.

Vision Transformer (ViT-H/14) achieved 88.55% Top-1 accuracy.

All benchmark tests verified under zero-hallucination grounded RAG verification.

Page 3 -- Citation Summary
Grounding accuracy: 99.4%
Citation latency: < 50ms
End of sample document.